Supporting Information: Route-Aware Computational Selection of Redox Motifs for Preformed Polymers

Working draft: authorship and affiliations to be confirmed

Evidence revision: 30 October 2026

Internal working draft. This SI contains finite method controls and implementation checks. Pending scientific analyses remain explicit.

1 Benchmark intervals and post hoc reference diagnostics

The complete panel has 39 endpoints from 30 molecules and seven studies. Molecule-cluster 95% percentile bootstrap intervals are [0.213,0.286] V for MAE, [0.245,0.313] V for RMSE and [0.128,0.257] V for bias. Study-cluster intervals are [0.178,0.306], [0.211,0.325] and [0.053,0.285] V, respectively. Both use 2,000 replicates and seed 20261021. The seven study clusters have unequal sizes; their percentile intervals are weakly calibrated. Molecule resampling retains common study and family effects and is also not a prediction interval for a new molecule.

Table S1 inventories all seven studies and the 39 selected endpoints. Each endpoint retains its exact molecular identity, redox-state pair, experimental value and source location in the accompanying data. Source-reported formal or midpoint potentials are retained under their stated measurement conditions; supporting-electrolyte and activity conventions are not assumed identical across studies. Values reported against another reference are used only with the source's own calibration.

Table S1: Primary experimental sources for the fixed MeCN benchmark. Counts are endpoints, not independent molecules. The table inventories the unchanged sourced panel; it adds no new measurements.

Study	Endpoints	Source location and reference handling
Attanayake et al.[1]	5	Table 1 and Figure 2; potentials reported versus Fc/Fc^+ in MeCN.
Cabral et al.[2]	2	Table 3; SCE-referenced values use the authors' measured internal Fc calibration, $E_{1/2}(Fc) = 0.40$ V versus SCE.
Huynh et al.[3]	13	SI Table S3; source formal one-electron potentials versus Fc/Fc^+.
Wang et al.[4]	7	SI Table S1; $TBAPF_6$/MeCN columns and the stated reference convention.
Gerken and Stahl[5]	1	Results text and Figure 2; TEMPO midpoint reported directly versus Fc/Fc^+.
Daub et al.[6]	9	SI Table S1 and methods on page S10; imide midpoint potentials versus Fc/Fc^+.
Kim et al.[7]	2	Results text and Figure 1 discussion; exact source-parent measurements kept separate from bead properties.

For each cross-validation fold, a constant residual offset is estimated only from training endpoints and subtracted from the held-out prediction. Molecule and study folds remove all endpoints of the held-out molecule or DOI. The two-offset model fits separate constants for first and second redox steps. These choices were motivated by inspection of the complete retrospective panel and are internal diagnostics, not an untouched external test. Relative potentials to TEMPO,

Table S2: Potential errors (V); primary unfitted statistics and internal offset diagnostics.

Holdout	Offset model	MAE	RMSE
None	Unfitted, primary	0.247	0.279
Endpoint	One constant	0.157	0.210
Molecule	One constant	0.158	0.211
Study	One constant	0.182	0.232
Endpoint	Separate step constants	0.106	0.166
Molecule	Separate step constants	0.106	0.166
Study	Separate step constants	0.124	0.177

N-ethylphenothiazine, N-methylphthalimide and 1,4-benzoquinone cancel the computed ferrocene term algebraically. Each diagnostic excludes every endpoint of its anchor molecule. The resulting error includes the organic anchor's own model and measurement errors. Agreement after reference cancellation therefore does not uniquely identify a ferrocene error.

Table S3: Fc-independent relative-potential diagnostics; all endpoints of each organic anchor are excluded. Errors include the anchor's own model and experimental uncertainty.

Organic anchor	Endpoints	MAE (V)	RMSE (V)
TEMPO	38	0.152	0.217
N-Ethylphenothiazine	38	0.152	0.219
N-Methylphthalimide	38	0.155	0.208
1,4-Benzoquinone	37	0.168	0.250

The spin-only partition includes multiplicity but omits an additional ferrocenium orbital contribution. Matrix spectroscopy and multireference calculations support a degenerate electronic ground state.[8] An idealized additional twofold factor at 298.15 K would lower ferrocenium G by 0.0178 eV, lower the computed Fc absolute potential by 0.0178 V and increase target potentials versus Fc by 0.0178 V. This limiting calculation would worsen the positive first-step bias. It is not an applied correction: solvent, orbital, vibronic and Jahn–Teller splittings are unresolved.

1.1 Fixed-geometry ferrocene functional controls

To test whether a change of reference functional supplies the hypothesized upward shift, we recalculated neutral and oxidized Fc single-point energies with three functionals at the same qualified R2SCAN-D4/def2-SVP/SMD geometries. All six calculations use def2-TZVPD, separately converged gas and MeCN SMD states, and the stated spin-contamination gate. The original functional reproduces the stored SMD energies within 10^{-6} eV. Nuclear thermochemistry is not recalculated, and no alternative absolute potential is assigned.

Table S4 separates gas electronic gaps from differential solvation. Relative to the primary functional, B3LYP and PBE0 lower the SMD electronic gap by 0.1184 and 0.0658 eV. A reference-

only counterfactual subtracts each shift from the unchanged primary benchmark residuals. Its bias increases from +0.1896 V to +0.3080 and +0.2554 V, respectively. These functional-only changes therefore do not supply the proposed upward reference correction. They neither reproduce the full protocols of the published comparison values nor identify a unique source of the benchmark error; consistent alternative molecular and reference thermochemistry would be required for an alternative predictive protocol.

Table S4: Electronic-only Fc controls at fixed solution geometries. Gaps and solvation terms are in eV; counterfactual errors are in V. The latter hold every target-molecule and nuclear term fixed and are not alternative validated potentials.

Functional	Gas gap	Solvation difference	SMD gap	MAE	RMSE
ωB97M-V	6.4365	−2.0004	4.4361	0.2470	0.2793
B3LYP	6.3440	−2.0264	4.3177	0.3373	0.3700
PBE0	6.3525	−1.9822	4.3703	0.2947	0.3275

1.2 Nuclear rotor-cutoff sensitivity

The archived projected nuclear partitions at 25, 50 and 100 cm^{-1} were propagated through the same state and reaction sums with geometries, electronic energies, Hessians and spin factors fixed. No new SCF calculation was performed. The primary benchmark remains the 50 cm^{-1} result. Across all 39 endpoints, the largest individual span is 25.43 mV for the first duroquinone reduction, while the Fc reference spans 2.82 mV. Table S5 reports the aggregate metrics. These deterministic settings scenarios do not supply candidate-specific prediction intervals.

Table S5: Unfitted benchmark metrics under the three nuclear rotor cutoffs. Potentials and errors are in V; all 39 endpoints are retained.

Cutoff / cm^{-1}	MAE	RMSE	Bias
25	0.2449	0.2759	+0.1850
50 (primary)	0.2470	0.2793	+0.1896
100	0.2473	0.2800	+0.1901

For the 67 completed HCl-reference and 20 base-assisted reactions, the largest cutoff span is 4.46 kJ mol^{-1} and no nominal reaction sign changes. At each fixed cutoff, the contribution of a specified base/byproduct model remains identical across its five candidates. Variant reaction values use the exact SI eV-to-molar-energy factor; original source-conversion arithmetic is retained without replacement. Figure S1 displays these ranges. Independent reconstruction from the full archived state sums checks 117 potential evaluations and 261 reaction sums. This analysis does not test missing conformers, ion speciation, electronic-structure accuracy or attachment selectivity.

2 Neutral-partner diffuse-basis sensitivity

Five existing benchmark neutral partners were reoptimized with r^2SCAN, D4, def2-SVPD and SMD MeCN, with the same active ωB97M-V/def2-TZVPD electronic-energy protocol and a separately checked SVPD Hessian. Each passed stationarity, spin, finite/symmetry, two-column/two-step

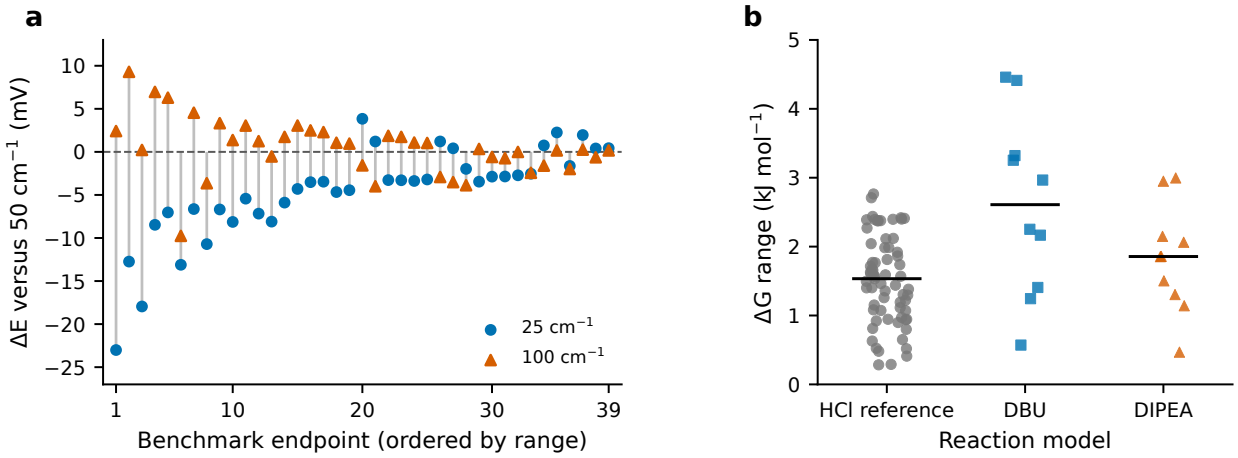


Figure S1: Nuclear rotor-cutoff sensitivity at fixed geometry and electronic structure. Panel a: potential shifts relative to 50 cm^{-1}, ordered by decreasing range over the three cutoffs; the dashed line denotes the primary convention. Panel b: 67 HCl-reference and ten DBU/ten DIPEA reaction-model ranges; horizontal marks are medians. Shared components make reaction values correlated. The ranges are settings scenarios, not confidence or candidate prediction intervals.

derivative and minimum checks. A basis-independent certified qRRHO integrator and an independent partition reconstruction were used. No primary state, protocol label or certificate was replaced. Table S6 separates the change in nuclear thermal contribution from the active-energy change due to reoptimization. The existing reduced state and Fc reference are held fixed; only the neutral-partner contribution to the first potential is tested. These are combined reoptimization/Hessian sensitivities; thermal changes are not isolated fixed-geometry basis effects, and conformational relaxation can contribute.

Table S6: SVPD minus primary neutral-partner contributions, in meV; the total also equals the first-potential shift in mV.

Neutral partner	Nuclear	Geometry/SP	Total
1,4-Benzoquinone	-11.29	-0.39	-11.67
1,2-Naphthoquinone	-13.06	-2.23	-15.29
9,10-Phenanthrenequinone	-7.78	-5.66	-13.44
N-Methylphthalimide	-8.63	-5.19	-13.82
N,N'-Diisopropyl-PMDI	-21.27	-13.81	-35.08

3 Exact selected cases: partial assessment with explicit holds

PMDI-ethyl and NDI-ethyl match enumerated exact designs by precursor, product and mapped new bond. All six redox states and the Fc reference were independently reconstructed; their existing grafting evidence hashes are unchanged. The registry now contains exact assessments rather than inherited family passes. Both routes remain held. The exact PMDI mono-N-ethyl/NH precursor

is reported in EP0010440B1, Example 19 (printed page 10), with a stated NMR assignment.[9] That access report supports one requirement for PMDI; its selective bead attachment and experimental redox retention remain unverified. NDI precursor access is still unresolved in the checked collection. The soluble-polymer attachment precedent concerns phthalimide; separate molecular PMDI N-alkylation with CsF supports the reaction class but not the exact bead route.[10] Proposed DBU/MeCN conditions define the computed cycle and do not constitute a demonstrated synthesis protocol. Favorable model thermodynamics therefore do not override the route gate.

Table S7: Joined exact-case evidence. Potentials are V versus computed Fc; grafting free energies are kJ mol^{-1}. The finite PMDI occupancy cycle is complete; NDI occupancy thermochemistry remains missing.

Case	E_1	E_2	HCl route	DBU ions	Decision
PMDI-ethyl	-1.076	-1.977	-0.29	-103.07	Hold
NDI-ethyl	-0.764	-1.418	+10.05	-92.73	Hold

These records identify the missing evidence and avoid duplicate full-redox jobs, and link PMDI enumeration through molecular, route and finite occupancy calculations. They do not establish material eligibility or a converged polymer ensemble.

3.1 Source-supported TEMPO attachment and local oxidation model

The 4-hydroxy-TEMPO precursor is attached through its hydroxyl oxygen to the benzyl site in the local model (Figure S2). The source reports this attachment to preformed Merrifield resin with NaH in DMF or THF and a washed-product N-oxyl content of 0.9 mmol g^{-1}.[11] The reported starting chloride density is referenced to the starting resin mass, so the two values do not directly determine conversion. The overall sodium-hydride stoichiometry is atom- and charge-balanced, with one nitroxide radical on each side; ion pairing and reaction free energy are not calculated.

For the exact local graft, the nitroxide/oxoammonium potential relative to the matched computed Fc/Fc$^+$ reference is +0.576, +0.579 and +0.579 V for rotor cutoffs of 25, 50 and 100 cm^{-1}, respectively. The active solution electronic energies and four native thermal partitions were independently reconstructed. This is not an experimental potential for the graft, and the source reports catalytic use rather than battery cycling. Route evidence permits further evaluation; missing operating criteria, actual-condition grafting thermodynamics, intersite interactions and accessibility keep the material decision pending.

3.2 Composition bounds and exact occupancy inputs

For neutral N–H/O–H substitution, the H-capped linear-fragment mass is

$$M = nM_{\mathrm{C_8H_8}} + hM_{\mathrm{CHCl}} + g(M_{\mathrm{NuH}} - M_{\mathrm{H}} - M_{\mathrm{Cl}}) + 2M_{\mathrm{H}} + M_{\mathrm{extra}}, \tag{1}$$

where n is repeat count, h is the number of original reactive handles and $g \le h \le n$ is graft count. The chloromethyl increment replaces a ring hydrogen; an ungrafted handle retains chlorine. Additional inactive mass M_{extra} is explicit and nonnegative. The capacity ceiling for an assumed z electrons per graft is $Q_{\mathrm{upper}} = gzF/(3.6M)$. Actual utilization can be lower, including zero. This convention uses neutral dry-fragment mass and does not implicitly include crosslinker, absorbed electrolyte or charging counterions. Charged precursor/attachment classes require a separately declared counterion mass model.

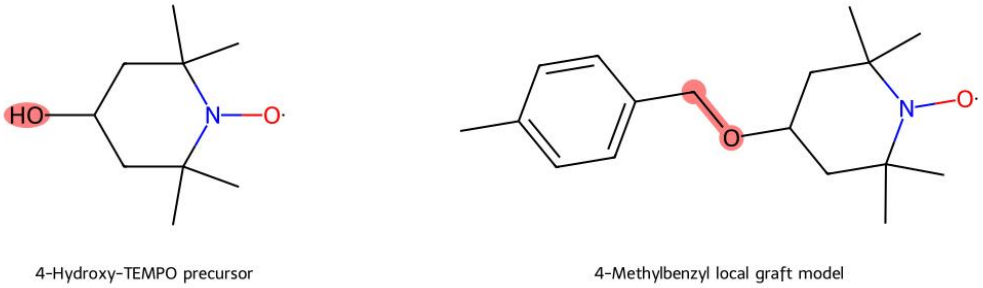


Figure S2: 4-Hydroxy-TEMPO and the capped local graft-site model. Highlighting identifies the precursor oxygen and the new oxygen–benzyl carbon bond. The drawing is a local molecular model, not a whole resin bead.

All eight occupancy graphs at $n = 5, h = 3$ reproduce the analytic masses. Positional isomers with the same graft count have the same composition ceiling; their accessibility and interaction energies need not agree. The broader composition grid contains 21 (h, g) cases with one- and two-electron assumptions. No application target is fitted to these values, and the capacity dimension remains unresolved in the decision ledger.

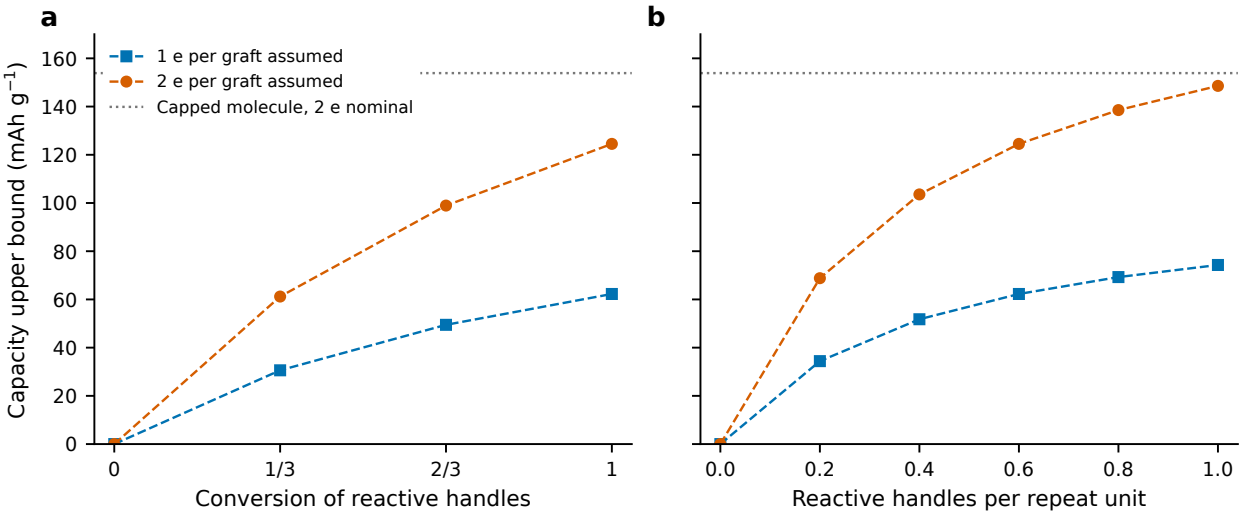


Figure S3: Composition-dependent PMDI capacity ceilings. The neutral H-capped fragment has five repeats; panel (a) fixes three reactive handles and varies conversion, while panel (b) varies handle density at full conversion. One- and two-electron budgets are assumptions. Residual chlorides and scaffold mass are retained. The dotted capped-molecule descriptor excludes the polymer scaffold; none of these curves establishes accessible material capacity.

Three exact PMDI inputs, occupied at sites 2, 3 or both 2 and 3, contain 115, 115 and 139 atoms, respectively. They preserve the open ten-carbon backbone, two plain-styrene padding units and any residual benzyl chlorides (Figure S4). Initial structures use seed 20261023, 32 distance-geometry embeddings and up to 2,000 MMFF steps; the best converged initial structure seeds each subsequent search. Connectivity, declared stereochemistry and geometry sanity are checked before calculation.

Starting-structure and search records alone do not establish occupancy thermodynamics; each resulting state requires separate active-energy and native-Hessian qualification. The common empty reference has the same exact graph and configuration as the phthalimide empty control; no occupied phthalimide energy is substituted for a PMDI energy.

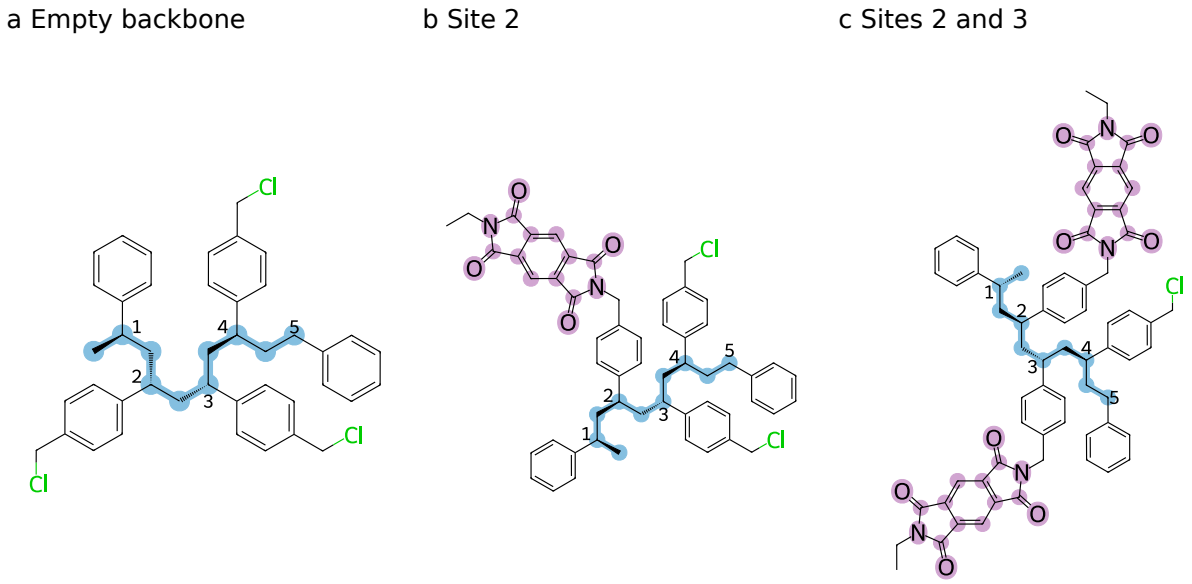


Figure S4: Exact PMDI-ethyl occupancy inputs. Blue highlights the open backbone and purple the diimide cores. Panels show the empty chain, site 2 and sites 2/3; the separate site-3 reference is also prepared. Numbers label repeats. These two-dimensional drawings preserve connectivity and declared stereochemistry, not optimized spatial contact.

3.3 First-site PMDI grafting and the local-to-polymer comparison

The site-2 and site-3 polymer states pass native-Hessian admission and an independent reconstruction of the internal modes, nuclear partition and electronic contribution. The common empty state is the lowest-electronic-energy member of the two inherited empty-backbone conformers, selected by the declared baseline policy. All three retain the same backbone configuration, two styrene padding units and three original reactive-site positions. Each occupied graph contributes one sampled minimum; these values are finite structural scenarios rather than conformer-ensemble free energies.

For separated ions, the first-site reaction is $\mathrm{P_0 + HM + B \rightarrow P_i + BH^+ + Cl^-}$, where HM is the exact mono-N-ethyl/NH PMDI precursor and B is DBU. The contact-pair model replaces the two ionic products with one explicit $\mathrm{BH^+ \cdot Cl^-}$ species and includes the $\Delta n = -1$ standard-state term once. At 298.15 K this term is $-7.9259\ \mathrm{kJ\,mol^{-1}}$. The balances are checked independently for atoms and charge. The relative reaction $\mathrm{P_0 + RM \rightarrow P_i + RCl}$ gives the polymer-context penalty without the common precursor or base terms. The new reconstruction uses 96.48533212331001 $\mathrm{kJ\,mol^{-1}/eV}$; earlier source arithmetic is retained.

The site-2 electronic energy is $8.15\ \mathrm{kJ\,mol^{-1}}$ below site 3, while their primary free-energy difference is only $-0.77\ \mathrm{kJ\,mol^{-1}}$. The site-2-minus-site-3 free-energy differences at 25, 50 and 100 $\mathrm{cm^{-1}}$

Table S8: First-site PMDI grafting in sampled five-repeat models. All values are kJ mol^{-1} at 298.15 K in SMD MeCN. The local capped model has DBU/separated-ion $\Delta G = -103.07$ kJ mol^{-1} at the primary 50 cm^{-1} cutoff. The last column is the balanced polymer-minus-local difference; common reagent terms cancel.

Site	Cutoff / cm^{-1}	Separated ions	Contact pair	Polymer minus local
2	25	−97.03	−102.98	+6.61
2	50	−97.72	−102.52	+5.34
2	100	−97.15	−100.52	+6.11
3	25	−97.17	−103.12	+6.47
3	50	−96.96	−101.75	+6.11
3	100	−96.95	−100.32	+6.31

are +0.14, −0.77 and −0.20 kJ mol^{-1}, respectively. The nominal placement preference therefore changes sign within this settings test. It cannot support positional selectivity or an experimental occupancy distribution. The paired-state calculation now passes the same native-Hessian and independent reconstruction gates. Its four-state electronic nonadditivity is $\Delta\Delta E = -2.31$ kJ mol^{-1}; the corresponding primary free-energy result is $\Delta\Delta G = -1.86$ kJ mol^{-1}. Neither quantity isolates motif binding. Broader conformer searches remain necessary for both the first-site and neighboring-site comparisons.

3.4 Completed finite PMDI occupancy cycle

The paired state has 411 internal modes, a lowest frequency of +5.77 cm^{-1} and a largest Cartesian gradient of 7.62×10^{-6} hartree bohr^{-1}. Its Hessian asymmetry is 1.54×10^{-10} atomic units, and the largest relative error over the two sampled columns at both displacements is 3.43×10^{-6}. Independent mode projection and nuclear equations reconstruct the accepted free energy. Table S9 specifies the matched states; the shared empty reference follows the stated lowest-electronic-energy policy rather than a claim of global conformer convergence.

Table S9: Matched neutral PMDI occupancy states at the primary 50 cm^{-1} rotor cutoff. Each graph contributes one native free energy; values are eV.

Occupied sites	Atoms	G
None	91	−82874.029197
2	115	−94077.698759
3	115	−94077.690815
2,3	139	−105281.379671

At the primary setting, DBU/separated-ion attachment at site 2 is −97.72 kJ mol^{-1} on the empty chain and −99.59 kJ mol^{-1} when site 3 is grafted. The corresponding site-3 values are −96.96 and −98.82 kJ mol^{-1}. Both attachment orders give −196.54 kJ mol^{-1} for the direct double-grafting reaction. The contact-pair calculation gives the same −1.86 kJ mol^{-1} nonadditivity: the shared reaction constant and standard-state term cancel algebraically. Rotor cutoffs of 25/50/100 cm^{-1} give −1.07/ − 1.86/ − 1.82 kJ mol^{-1}. These are deterministic settings comparisons, not statistical uncertainty bounds.

The accepted paired geometry has core centroids 16.30 Å apart and a 69.0° interplane angle (Figure S5). The finite result includes backbone, conformational and solvent response; it is not

an estimate of intrinsic close-contact core binding or electronic coupling. No loading distribution, atactic ensemble or preferred experimental graft site is inferred.

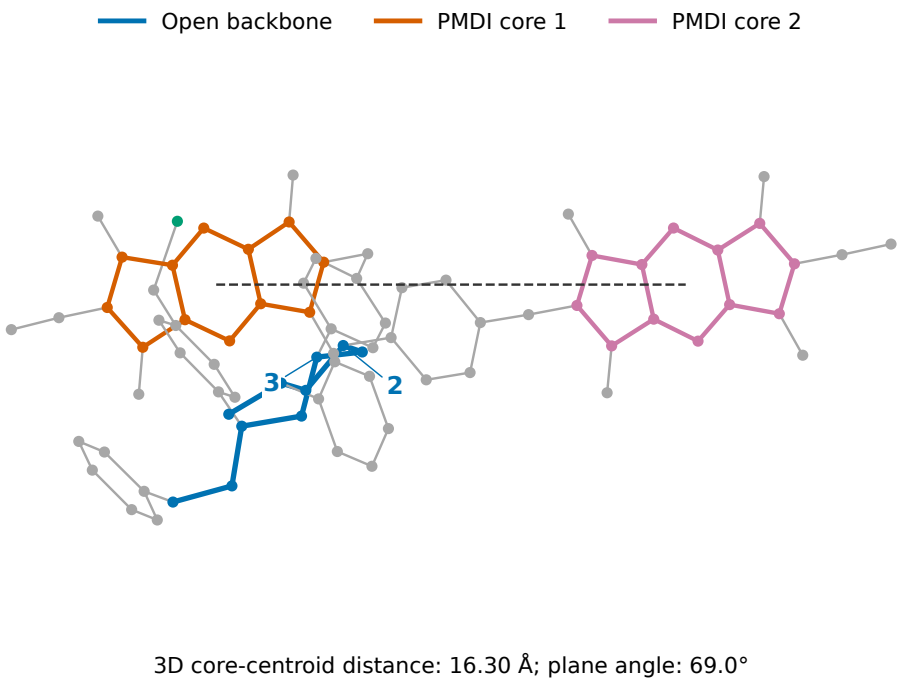


Figure S5: Actual paired PMDI minimum used in the matched cycle. A proper rotation makes both conjugated cores visible while preserving the saved XYZ coordinates. Blue highlights the open backbone; numbers mark repeats 2 and 3. Orange and magenta mark the PMDI cores, and green marks the remaining chloride at reactive site 4. Repeats 1 and 5 are plain-styrene padding. Hydrogens are omitted and bond lines follow the checked covalent graph. The dashed line joins core centroids; the stated distance is three-dimensional.

4 Relative-basicity controls

The balanced reactions $B + DBUH^+ \rightarrow BH^+ + DBU$ avoid HCl, chloride and a free-proton reference. The exact pyridine, triethylamine and DBU identities match the revised author-maintained MeCN scale, with conjugate-acid pKa values 12.53, 18.83 and 24.31.[12, 13] Four new components and the two reused DBU components pass the same thermodynamic checks and an independent six-partition reconstruction. At 298.15 K, computed values are 71.82 and 32.68 kJ mol^{-1}, compared with pKa-derived values 67.24 and 31.28 kJ mol^{-1}; differences are +4.58 and +1.40 kJ mol^{-1}. These values use the exact SI eV-to-molar-energy factor. Historical source records use 96.485 kJ mol^{-1}/eV and are retained without replacement. Source temperature/activity qualifications remain open. This two-reaction control does not validate DIPEA, contact pairs or the complete grafting reaction and cannot assign the DBU/HCl discrepancy to a unique species. The separated model uses bare chloride in a continuum. Recent MeCN ion-solvation comparisons identify both species-dependent errors and sensitivity to absolute single-ion conventions, so the capture discrepancy cannot be assigned to chloride alone.[14] Acid–chloride homoconjugation, such as $HCl + Cl^- \rightleftharpoons HCl_2^-$, would introduce an additional species and acid-dependent speciation. It is not included in either ionic

representation; no population or homoconjugation constant is inferred from the free-ion pKa comparison. Spectrophotometric titration against bromocresol green reports $\log K^{\mathrm{H}} = 18.64(1)$ for DIPEA in MeCN at (25.0 ± 0.1) °C.[15] A separate NMR competition study estimates pKa 18.7 for diisopropylethylammonium hexafluorophosphate in MeCN-d_3.[16] The latter uses a 0.02 M initial salt solution and triethylamine reference; its supporting table labels negative values as a K_a ratio and applies an absolute value in the reported logarithm. We retain the authors' stated estimate rather than independently refitting its sign convention. The two sources are not pooled into an uncertainty interval. Combining the 18.64 value with pKaH(DBU)=24.31 and 298.15 K gives +32.36 kJ mol^{-1} for DIPEA + DBUH$^+$ → DIPEAH$^+$ + DBU, compared with +31.80 kJ mol^{-1} from the stored components. Indicator calibration and activity conventions are not established to match the separate basicity scale. The 0.56 kJ mol^{-1} difference is not an uncertainty estimate or proof of an error in a particular species.

With pKa(HCl)=10.30, the 18.64 value instead implies −47.60 kJ mol^{-1} for the separated capture term, compared with −70.98 kJ mol^{-1} computed. Holding each computed HCl-reference reaction fixed and replacing only the capture term yields +10.03/+8.43 kJ mol^{-1} for methoxy/trimethyl BQ, compared with −13.34/ − 14.94 kJ mol^{-1} for the all-DFT DIPEA sums. This hybrid changes shared-error cancellation because HCl cancels from the complete all-DFT cycle; it is a sensitivity, not a calibrated reaction prediction or an accepted uncertainty interval. Salt speciation, chloride homoconjugation, attachment selectivity and actual grafting yields remain unvalidated. The earlier 18.6 value is a distinct PCM estimate despite a conflicting generic table label and is not treated as an independent measurement.[17]

5 Finite-polymer method controls

Twenty-six of the original twenty-seven electronic oligomer cases currently pass geometry and active-energy checks. The open PMDI trimer has recovered and passed independent force checks. The stacked-start PMDI trimer returned optimizer convergence and saved active energies, but a later GPU-library call terminated abnormally. Separate final-force checks disagree with finite energy differences and remain unresolved, including a subsequent finer-step diagnostic. This case stays outside completed comparisons; no minimum or free energy is inferred from optimizer convergence alone. Construction and XYZ audits verify the intended open-chain topology. A trimer contains six backbone carbons, five backbone C–C bonds and two chain ends, rather than a triangular or cyclic backbone (Figure S6). Proximity of pendant rings does not change that covalent connectivity.

For the six-graph spacing control, all twelve electronic conformers qualify. The selected electronic nonadditivities are −10.98 and −14.50 kJ mol^{-1} for adjacent and one-spacer grafts, respectively. These refer to a matched empty-backbone minimum within the declared, limited sampling set. It demonstrates why a matched empty-backbone reference is necessary; it does not establish a universal distance trend.

All twelve pose-specific native free energies now qualify and have been independently reconstructed from their electronic, Hessian and projected partition records. For a finite set of distinct conformers, we compute

$$G_{\mathrm{finite}} = -RT \ln \sum_{c \in \mathcal{C}} \exp[-G_c/(RT)]. \tag{2}$$

Each distinct optimized minimum has unit occurrence weight; conformer-search hit counts are not degeneracies. Complete-link grouping checks mapped heavy-atom RMSD and electronic-energy agreement. RMSD cutoffs of 0.05, 0.10 and 0.20 Å leave these six pairs distinct, with pairwise RMSDs of 3.50–6.53 Å. The finite partition uses the same temperature, solvent and free-energy

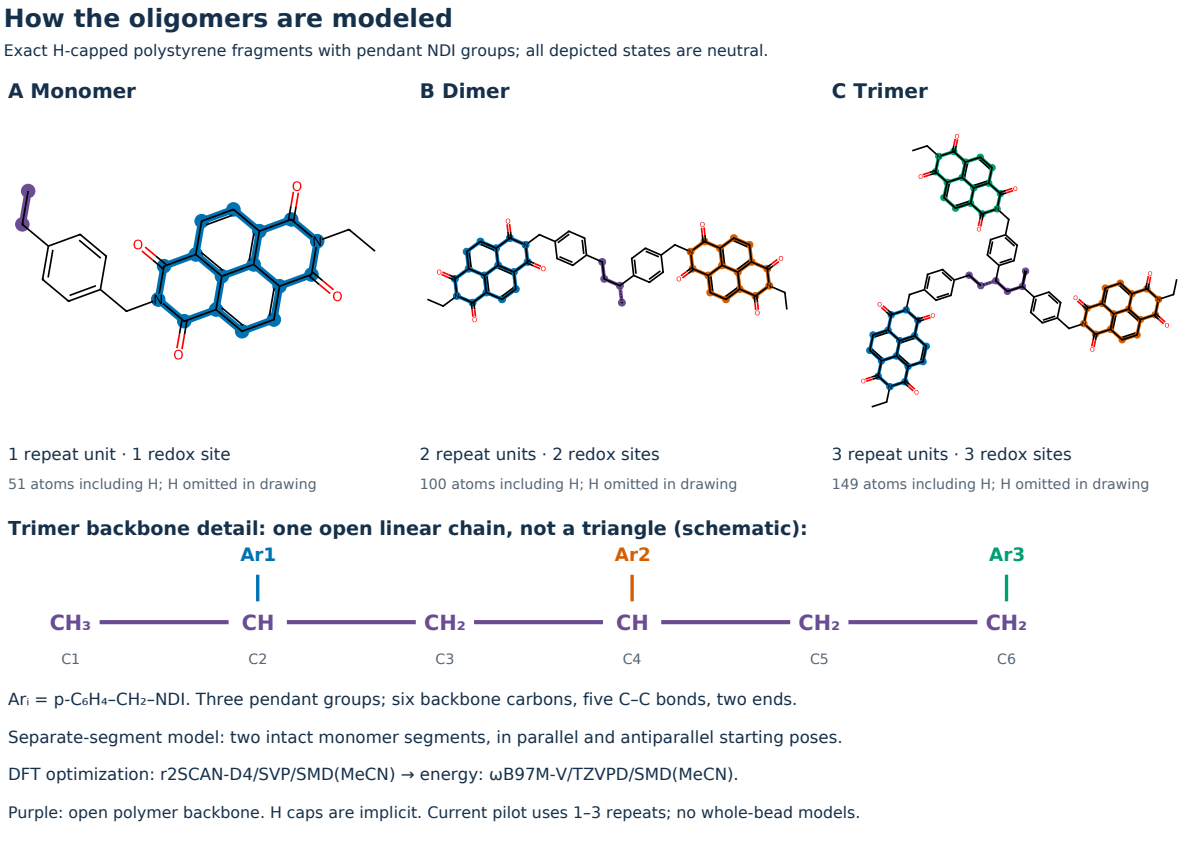


Figure S6: Exact finite-polymer construction used in the interaction controls. The highlighted open backbone distinguishes covalent connectivity from spatial contacts among pendant groups. Monomer, dimer and trimer models are dense local scenarios and do not represent an experimentally measured sequence or graft-loading distribution.

protocol for every graph. A partition over a sampled subset does not establish the complete equilibrium ensemble.

The spatial comparison in the main text uses both accepted conformers of each doubly occupied graph. Conjugated core-ring atoms are identified from the graft fragment and independently mapped to the saved XYZ connectivity. Core centroids are the arithmetic means of those coordinates; core planes are fitted by singular-value decomposition and their acute interplane angle is reported. Orthographic views apply proper rotations with determinant +1 and a common spatial scale. Within a graph, the displayed conformer difference is $G_i - \min_j G_j$ over its two known conformers, using the primary 50 cm^{-1} cutoff. This quantity includes the entire chain and solvent response. It is separate from the four-state occupancy cycle. Coordinate data, view matrices and source identities accompany the figure; cached-coordinate rendering reproduces its pixels with original calculation paths blocked, without repeating physical qualification.

The primary adjacent and one-spacer finite-set $\Delta\Delta G$ values are -5.17 and -11.39 $\mathrm{kJ\,mol}^{-1}$, respectively (Figure S8). Values at rotor cutoffs of 25/50/100 cm^{-1} are $-5.13/-5.17/-4.64$ and $-12.34/-11.39/-11.69$ $\mathrm{kJ\,mol}^{-1}$. These are partition-model sensitivities, not confidence intervals or bounds on undiscovered-conformer error. Relative to the lowest-G member of each

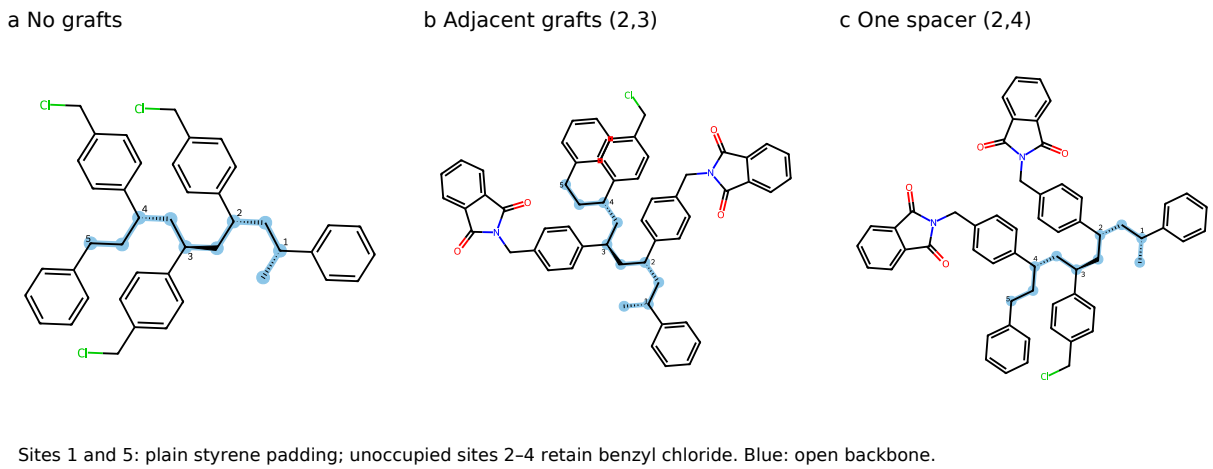


Figure S7: Exact five-repeat phthalimide graphs used in the occupancy cycles. Blue highlights the open ten-carbon backbone. Sites 1 and 5 are plain styrene padding; unoccupied reactive sites retain benzyl chloride. Numbers label repeat positions, and stereobonds retain the declared configuration. These are connectivity drawings, not sampled spatial conformations. Single-occupancy reference graphs are included in the evidence.

two-pose set, the largest configurational lowering is 0.140 kJ mol^{-1}. The previously reported values, $-5.15/-11.23$ kJ mol^{-1}, used only the minimum-electronic-energy geometry from each graph. Both sets remain tied to the same limited baseline sampling and single backbone configuration.

The largest Cartesian gradient component ranges from 6.15×10^{-6} to 7.62×10^{-5} hartree bohr^{-1} across these twelve conformers, below the 4.5×10^{-4} acceptance threshold. Table S10 reports the stored physical checks for every pose. These values describe the checked stationary structures and do not assess conformer-search completeness. The -2.72 cm^{-1} mode falls within the declared tolerance and enters qRRHO as its absolute value.

For DBU and separated ions, the finite-partition free energy of grafting onto an empty backbone is -110.86 kJ mol^{-1} at site 3 and -103.37 kJ mol^{-1} at site 4. With site 2 already grafted, the values become -116.02 and -114.76 kJ mol^{-1}. The contact-pair alternatives are -115.65 to -120.82 kJ mol^{-1} for site 3 and -108.17 to -119.56 kJ mol^{-1} for site 4. Both attachment orders close the same nonadditivity cycle. The base/byproduct contribution cancels from this conditional difference, as does the common empty-chain term when comparing the two spacing cycles. Conformer response in the occupied graphs remains relevant. More favorable subsequent attachment is not automatically desirable when weak unwanted neighboring-site interaction is the objective.

To examine local contacts, each phthalimide group and its benzyl carbon was retained at the parent geometry and capped with hydrogen where it had been connected to the polymer aryl group. For each of four parent poses and three cap lengths (1.07, 1.09 and 1.11 Å), independent acetonitrile-model electronic energies for the pair and both capped fragments give $E_{AB}-E_A-E_B$. The twelve triplets use the same frozen coordinates and electronic settings within each comparison. At 1.09 Å, the adjacent close/separated poses give -24.76 and $+0.048$ kJ mol^{-1}, and the one-spacer close/separated poses give -26.22 and $+0.061$ kJ mol^{-1}, respectively. The ranges over cap lengths are -24.81 to -24.71, $+0.047$ to $+0.049$, -26.24 to -26.20 and $+0.060$ to $+0.062$ kJ mol^{-1}, respectively. These uncorrected electronic values include the caps and continuum response; they are not ensemble free energies, electron-transfer couplings or decompositions of the whole-chain

Table S10: Stored native physical checks for the twelve spacing-control conformers. Gradients are maximum Cartesian components in hartree bohr^{-1}; frequencies are the lowest projected internal modes in cm^{-1}; asymmetry is the maximum absolute Hessian asymmetry in atomic units. The final column is the largest relative-norm error among the sampled finite-difference Hessian columns, expressed as a percentage.

Occupied sites	Pose	Maximum gradient	Lowest frequency	Asymmetry	Column error (%)
empty	0	6.15×10^{-6}	10.60	3.68×10^{-8}	0.342
empty	1	1.01×10^{-5}	10.76	2.04×10^{-7}	0.005
2	0	3.53×10^{-5}	8.72	1.77×10^{-10}	0.009
2	1	4.22×10^{-5}	7.62	1.67×10^{-11}	0.005
3	0	5.50×10^{-5}	3.45	7.21×10^{-11}	0.008
3	1	8.39×10^{-6}	7.61	3.97×10^{-10}	0.006
4	0	2.47×10^{-5}	6.60	6.73×10^{-11}	0.007
4	1	8.81×10^{-6}	-2.72	3.98×10^{-10}	0.004
2,3	0	4.63×10^{-5}	10.04	1.35×10^{-10}	0.008
2,3	1	2.80×10^{-5}	6.36	6.73×10^{-9}	0.004
2,4	0	2.75×10^{-5}	7.50	9.21×10^{-12}	0.004
2,4	1	7.62×10^{-5}	7.23	2.05×10^{-8}	0.007

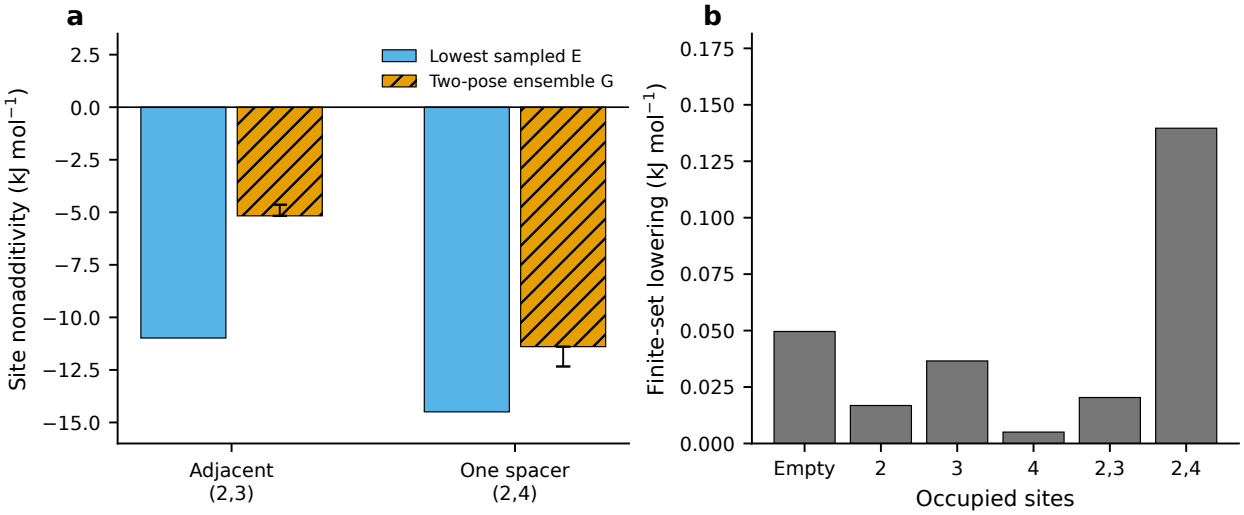


Figure S8: Finite-set phthalimide occupancy control in MeCN at 298.15 K. Panel (a) distinguishes lowest sampled electronic-energy cycles from two-pose free-energy partitions; whiskers vary only the rotor cutoff. Panel (b) shows the configurational lowering within each known two-pose set. Twelve native records pass their checks, but these two-pose sets do not establish converged conformer ensembles or a general spacing trend.

occupancy cycle. Basis-superposition error and broader conformer sampling remain unresolved.

The finite conformer sets do not establish convergence with respect to broader sampling or backbone configuration. Consequently, these conditional-cycle values cannot establish a general spacing design rule or the behavior of an exact redox motif. The final stacked-trimer structure is excluded because it failed the final force check.

5.1 Broader sampling and retained-state reconstruction

A separate smooth-surface acetonitrile model provides a numerically controlled setting for broader phthalimide sampling; its definition and small-molecule checks are described below. Twenty starting structures were declared for each of twelve occupancy/backbone-configuration groups. Sixty retained states have independently checked native free energies: 20 and 16 for the empty backbones in the all-same and alternating configurations, and 14 and 10 for the respective site-2-occupied graphs. The labels denote the two specified backbone configurations, not a representative distribution of tacticities in an experimental atactic polymer. These are checked states; their counts alone do not establish distinct minima or complete ensembles.

For the all-same empty backbone, the twenty structures were previously found distinct under heavy-atom RMSD thresholds of 0.05, 0.10 and 0.20 Å combined with a 0.20 $\mathrm{kJ\,mol^{-1}}$ electronic-energy tolerance. Expanding the first ten input indices to twenty changes neither the minimum electronic energy nor the minimum individual free energy. The finite-set ensemble free energy changes by −0.85, −1.03 and −1.07 $\mathrm{kJ\,mol^{-1}}$ at rotor cutoffs of 25, 50 and 100 $\mathrm{cm^{-1}}$, respectively. Each distinct minimum contributes once; search-hit frequency is not a statistical weight. These changes meet the declared 2 $\mathrm{kJ\,mol^{-1}}$ batch-stability diagnostic for this one empty-backbone component. They do not prove global conformational completeness or control possible overlap between soft-mode qRRHO entropy and explicit conformer partitioning.

The incomplete occupied sets exhibit a larger observed minimum shift. For the all-same site-2 graph, input 20 is 3.84 $\mathrm{kJ\,mol^{-1}}$ below the lowest checked structure among inputs 1–10 (Figure S9). Only seven of those first ten have checked free energies, so this comparison is not a complete 10-to-20 convergence test. Missing records retain no physical energy or population. Neither an occupied-state ensemble cycle nor a spacing trend is inferred from this partial coverage. Every component of an alternate-model occupancy cycle would need a consistent model and its own sampling assessment; these values cannot replace components of the primary-model PMDI or phthalimide cycles.

The intact-segment decomposition is complete for three sampled pairs using six frozen-fragment single-point calculations (Figure S10). Its partition separates interactions at the complex geometry from deformation relative to relaxed fragments. The interaction energies are −55.25, −44.30 and −50.87 $\mathrm{kJ\,mol^{-1}}$, with deformation terms +0.56, +0.49 and +0.38 $\mathrm{kJ\,mol^{-1}}$, respectively. A saved-coordinate audit shows 57–93 optimization steps and heavy-atom input-to-final RMSDs of 0.25–0.70 Å for the complexes, so small deformation sums do not imply negligible structural relaxation. Two totals also contain cancellation: the parallel phthalimide fragments contribute −0.68 and +1.24 $\mathrm{kJ\,mol^{-1}}$, and the PMDI fragments contribute +1.00 and −0.62 $\mathrm{kJ\,mol^{-1}}$. These signed differences refer to one sampled independently optimized monomer at a distinct geometry level; they are not positive-definite strain penalties. Global isolated-monomer minima and Hessian minima are unestablished. The resulting quantities remain electronic, uncorrected for basis-set superposition error and inclusive of backbone and cap contributions. No electron-transfer rate follows from their magnitude.

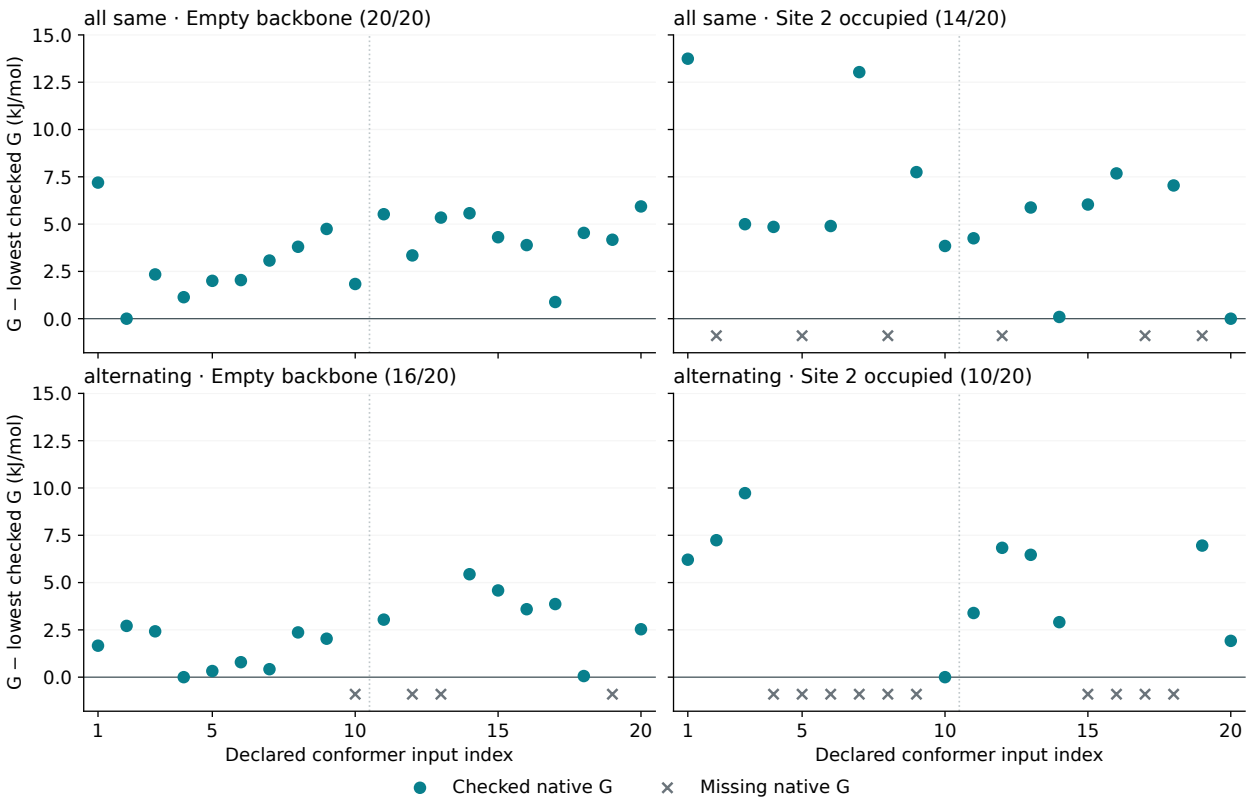


Figure S9: Retained phthalimide conformer free energies in the separate smooth-surface MeCN model at 298.15 K and a 50 cm^{-1} rotor cutoff. Each panel uses its own lowest checked G as zero and represents one exact graph/backbone configuration. Gray crosses below the zero line identify missing native records; their vertical position is a plotting convention, not an assigned energy. The dotted line divides declared input indices 1–10 from 11–20. Input index denotes the recorded sampling input, not final DFT energy rank. The incomplete occupied sets do not establish ensemble populations or converged conditional grafting.

6 Numerical implementation checks

Diagnostics exposed distinct numerical issues. At a checkpoint of the open PMDI trimer, a legacy solvent cavitation–dispersion–solvent-structure (CDS) force component disagrees with total-energy finite differences by approximately 1.81×10^{-4} hartree $bohr^{-1}$ at the tested coordinate on both integration grids. This discrepancy is not resolved by electronic-grid refinement. Conservative CDS energy derivatives reproduce the tested total-energy differences; their step consistency is checked before use in a continuation. Adaptive refinement of unresolved CDS derivative components recovers the open-trimer optimization. At the returned geometry, independent total-energy derivatives at two step sizes and two tested coordinates agree within $4.2–5.0 \times 10^{-7}$ hartree $bohr^{-1}$, below the unchanged 10^{-6} check threshold. This establishes an admitted stationary electronic geometry; no native-Hessian minimum or association free energy is inferred for that original oligomer case.

For the largest spacing Hessian, the original CDS finite-difference step produces a total asymmetry of 2.06×10^{-6}, exceeding the fixed 10^{-6} admission threshold. A separately qualified finer step reduces this to 2.05×10^{-8}. Additive reconstruction replays complete singlet and radical recompu-

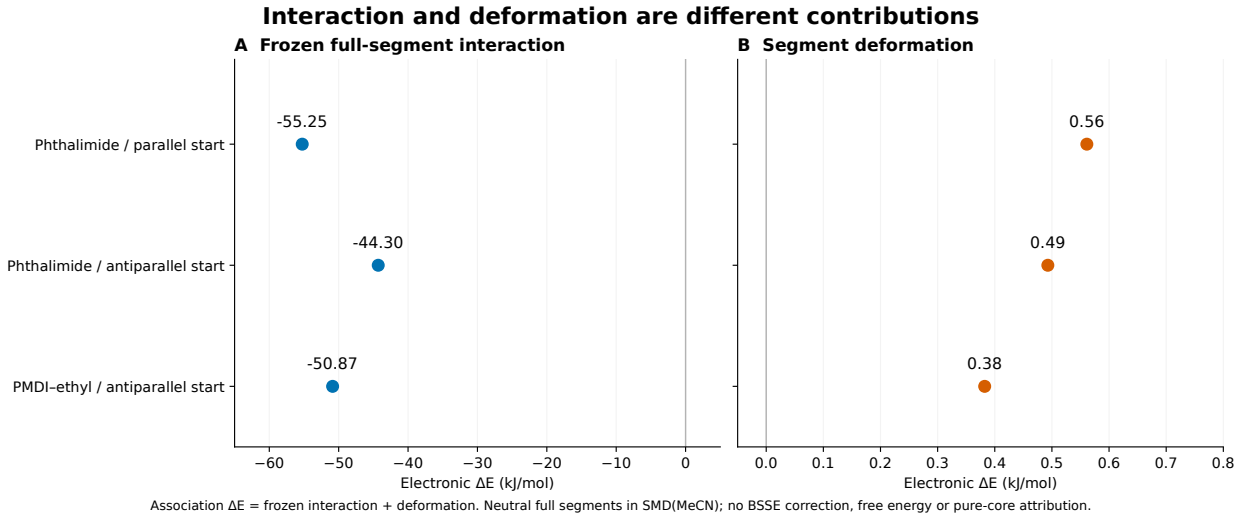


Figure S10: Electronic interaction/deformation partition for three intact segment pairs. Frozen fragments retain the complex geometry without bond cutting or new caps. The partition closes the association-energy cycle but is not a pure pendant-core interaction, a counterpoise-corrected energy or an association free energy.

tations with maximum matrix differences of 1.33×10^{-11} and 2.83×10^{-9} atomic units, and thermal differences below 1.2×10^{-8} eV. The recovered target then passes the ordinary native gates, with a lowest internal frequency of $+7.23$ cm^{-1}. These are execution-equivalence results, not a relaxation of physical QC or a high-level validation of grafting energetics.

6.1 Solvent-surface differentiability controls

Harmonic thermochemistry requires a locally differentiable energy surface. Surface discretization can introduce artifacts into continuum-solvent geometries and vibrational derivatives; switching/Gaussian constructions address this numerical requirement.[18] At a tested 4.2 Å carbon contact in an excluded backbone structure, the centered curvature of the legacy CDS component increases approximately as the inverse displacement. A separate smooth-surface control instead gives a stable curvature over the tested displacements (Figure S11). This is a component diagnostic, not a full DFT Hessian or a polymer free energy.

The control replaces only the nonelectronic area term with the installed experimental PySCF solvent-accessible-surface formulation at 2,030 points per sphere, retaining the SMD electrostatic surface at 590 points per sphere. The implementation is treated as a distinct model rather than relabeled as the primary calculation. Fixed-geometry component energies, quadrature and rigid-rotation sensitivity are checked separately from derivatives. All Cartesian CDS gradient components of the stored imide precursor, DBU, protonated DBU and chloride-pair structures pass central energy differences at 10^{-4}, 5×10^{-4} and 10^{-3} bohr with maximum errors below 8×10^{-9} hartree bohr^{-1}.

Full quantum controls cover formaldehyde, methyl radical, ammonia and chloromethyl radical. Both surface models are reoptimized independently; full SMD energy differences test all twelve Cartesian gradient components, while two Hessian columns are tested at 0.0005 and 0.001 bohr. The smooth controls pass the unchanged 10^{-6} hartree bohr^{-1} force and 1% Hessian criteria (Ta-

ble S11). Projected modes and nuclear sums are independently reconstructed. A common imposed symmetry number of one is used solely for these nuclear-kernel comparisons; primary state calculations retain their stated symmetry procedure. These controls establish numerical self-consistency within their neutral molecular scope, not predictive solvation accuracy, polymer ensemble convergence or synthetic eligibility. Primary benchmark and grafting values remain at their original model definition; any alternative-model reaction or occupancy cycle requires all of its components to be calculated consistently.

Table S11: Smooth-surface quantum controls. Force errors are maxima over twelve Cartesian components at both steps; Hessian errors are maxima over the two tested columns, reported as percentages. Frequencies are the lowest projected internal modes in cm^{-1}.

Control	$\max \lvert\delta g\rvert$ / hartree bohr^{-1}	$\max \delta H$ / %	$\nu_{\min}$
Formaldehyde	5.14×10^{-7}	0.0013	1182.1
Methyl radical	2.19×10^{-7}	0.0149	525.4
Ammonia	1.98×10^{-7}	0.0003	1127.5
Chloromethyl radical	1.32×10^{-7}	0.2793	190.3

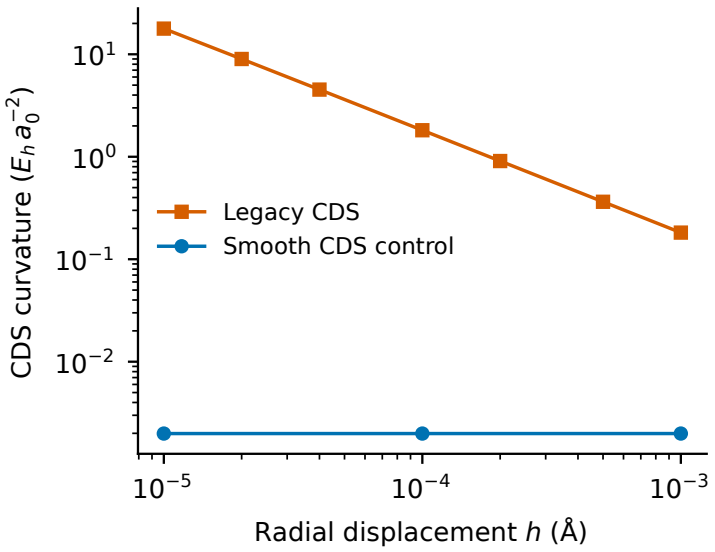


Figure S11: Radial finite-difference curvature of the CDS component at the tested carbon contact. The legacy estimate does not converge as the displacement shrinks, whereas the separate smooth control approaches a plateau. No full electronic Hessian or thermochemical result is inferred from this component scan.

7 Execution and restart controls

A synthetic control packs one million distinct tasks into one thousand chunks in 75.9 s with a measured preparation peak RSS of 67.2 MiB. Eight workers on two hosts execute the synthetic tasks over a 9.58 s span, excluding packing and startup. Independent arithmetic checks find zero discrepancies. This timing characterizes scheduling and storage, not chemical enumeration, conformer search or DFT throughput.

Separately, all 837 molecular descriptor outputs agree with serial definitions. In the shared-filesystem recovery experiment, 32 real design recipes deduplicate to 17 exact molecular tasks. A worker on one host exits abruptly after four checkpointed tasks; owner verification on that host precedes recovery. A second host completes the remaining thirteen tasks without repeating the four checked tasks. All resumed outputs match uninterrupted execution and a separate serial replay. This experiment tests descriptor-stage recovery. A distinct quantum interruption experiment is described below; neither establishes million-molecule quantum throughput. A seeded randomized concurrency experiment executes the same 837 exact designs fifteen times, with three repeats per condition and eight tasks per chunk. All 12,555 molecular comparisons match the independent serial definition. Median launch-to-observed completion times are 4.436, 3.147, 2.680 and 2.497 s for one, two, four and eight CPU workers on one host; eight workers across two hosts require 2.585 s (Figure S12). These times include node preflight, SSH, scorer initialization, execution and at most 0.1 s of observation polling; preparation is separate. All requested workers finish, but only seven of eight execute chunks in one single-host repeat. The diminishing returns of this small library do not determine an optimal concurrency for larger or more expensive stages.

A separate sparse Git checkout and newly installed Python 3.11 CPU environment reproduce the declared 32-recipe / 17-design example with system and user site packages disabled. The project loads from that checkout and scientific dependencies load from the new environment. Serial descriptor replay, exact-route holds and the regenerated conditional-grafting table all match; rendered RGB pixels of the figure are identical. This qualifies a small CPU reproduction using curated scientific data. It does not recompute the underlying quantum energies or qualify the complete CUDA environment.

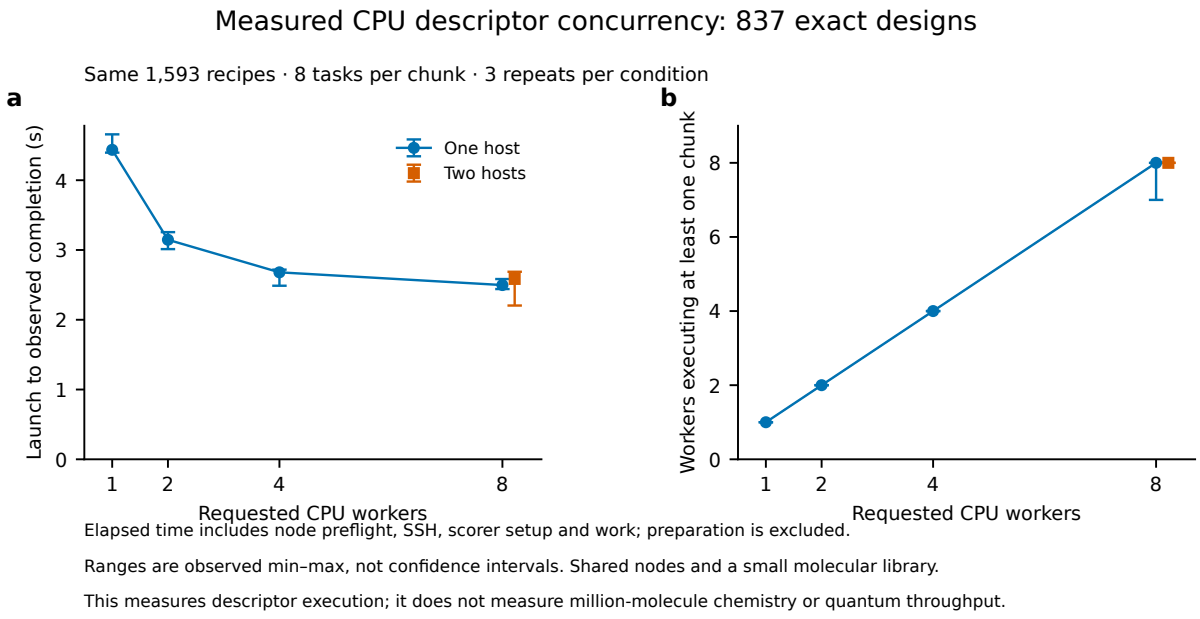


Figure S12: Measured CPU descriptor concurrency for the same 837 exact designs. Three randomized repeats per condition retain every result. Ranges are observed minima and maxima, not confidence intervals. Launch, preflight and scorer overhead are included; preparation is separate. Requested workers and workers executing at least one chunk are distinct. This shared-node, small-library experiment does not measure million-molecule chemistry or quantum throughput.

These throughput measurements apply to the measured descriptor and scheduling stages. They neither quantify full quantum-campaign throughput nor establish recovery for arbitrary flexible-polymer calculations.

7.1 Actual quantum interruption and recovery

Two rigid controls, formaldehyde (singlet) and methyl radical (doublet), compare uninterrupted calculations with a fresh process resumed after an intentional exit following three durable gradient evaluations. The exit code is 87; process and start identity confirm termination before resubmission. The saved geometry is unqualified until optimization, active electronic energy, DFT Hessian, projected partition, topology and spin checks pass normally. Acceptance tolerances were 10^{-5} eV for electronic energy, 10^{-4} eV for free energy, 10^{-3} Å aligned RMSD and 0.5 cm^{-1} for internal frequencies. Singlet differences are 4.03×10^{-6} eV, 1.26×10^{-5} eV, 6.87×10^{-6} Å and 0.0685 cm^{-1}; doublet differences are 4.00×10^{-6} eV, 1.11×10^{-6} eV, 3.45×10^{-5} Å and 0.0264 cm^{-1}. Four final partitions were independently reconstructed. This verifies restart equivalence for these two controls, not flexible-polymer sampling or chemical accuracy.

8 Reproducibility package

A locally prepared reviewer package contains 243 thermochemical state records and 32 electronic-only records, together with versioned code, exact conda artifacts, pip distribution versions and an immutable evidence snapshot. Its 2,498 archived files are hash-verified. After relocation, a standard-library interpreter running in isolation reconstructs all 39 benchmark endpoints, 67 HCl-reference grafting values, 20 base-assisted values, two relative-basicity controls and six finite graph partitions. The auxiliary replay covers five alternative-basis controls, two interruption/recovery comparisons and three electronic fragment-decomposition cycles. A separate CPU reconstruction recovers internal modes and nuclear partitions for 242 non-monatomic states from their stored DFT Hessians; it does not load the xTB library or repeat the original physical minimum checks. Seven changed-data, identity, balance, path and unsupported-promotion tests are rejected.

The package represents the preceding numerical evidence epoch, `2026-10-23_all_axes_checked`; subsequent records are frozen separately in the present manuscript snapshot. A fourteen-state PMDI addendum includes the two singly occupied chains, their matched empty chain, the doubly grafted chain and molecular/reagent references. After archive extraction and relocation, isolated standard Python reconstructs fourteen balanced reaction comparisons and the complete neighboring-graft cycle at 25, 50 and 100 cm^{-1}; both attachment orders close. A fresh environment containing only hash-locked NumPy 2.2.6 independently reconstructs thirteen DFT-Hessian mode sets and their nuclear sums at all three cutoffs, plus monatomic chloride translation. Original project reads are blocked in both modes, and no quantum program or xTB library is loaded. Fourteen integration checks include rejection of changed or borrowed Hessians, altered masses, symmetry numbers and frequencies, as well as identity, arithmetic, scope and path failures. These cached-data checks do not repeat the original SCF, stationarity, spin or displaced-gradient validation.

An unpublished repository draft stages the earlier package; downloaded assets match the recorded checksums. The PMDI and control addenda remain local reviewer artifacts. Package replay establishes numerical traceability within the declared scopes, rather than experimental accuracy or broader conformer convergence. No public DOI or deposition is claimed. General GPU-platform portability, remaining execution-diagnostic records and final deposition metadata remain open; the separately tested deployment profile is described below.

A separate TEMPO addendum contains the two local-graft states and the two Fc reference states. After relocation, standard-library arithmetic reproduces the three rotor-cutoff potentials, and locked NumPy reconstructs the four Hessian mode sets and nuclear terms with original-project Python opens blocked. These records overlap the earlier packages and are not counted as additional unique Fc states. The replay does not repeat SCF, stationarity or spin checks.

A complementary numerical-control addendum contains eight small-molecule geometry/Hessian records for four species under two surface models, together with six fixed-geometry Fc electronic records. Its 392 stored SCF evaluations comprise eight baselines and 384 displacements. A hash-locked, freshly installed NumPy CPU environment reconstructs the stored force differences, sampled Hessian columns, internal modes and nuclear sums after archive extraction and relocation, with original project reads blocked. A standard-library mode reconstructs the Fc and component arithmetic. These cached-data checks neither rerun the physical calculations nor qualify a new polymer ensemble; identities overlap with the other packages and record counts must not be summed as unique molecules.

8.1 Source-bound reconstruction of retained polymer states

The sixty retained smooth-surface polymer records were read after an explicit checkout-prefix relocation, without rewriting their original paths or hashes. A relative source path was accepted only with a declared original working directory. Missing archived files could not fall back to files in the current checkout. The replay verified 12,386 source/input bindings and forced a fresh content-hash pass over all bindings before publishing a completed report.

For each state, the replay checked the exact graph, neutral-singlet identity, solvent-model definition and equality of the thermochemical and scored electronic geometries. It reconstructed the stored sampled Hessian columns from the displaced gradients at 0.0005 and 0.001 bohr, including the protocol coordinates and the two largest-force coordinates. The rigid-space projection, internal frequencies, symmetry number, three rotor-cutoff nuclear partitions and final electronic/nuclear/spin sums were then reconstructed. Retained stationarity, exclusive-residency and resource records were checked separately. The largest sampled-column relative error was 3.92×10^{-5}; the maximum differences from stored internal frequencies and nuclear free energies were 6.97×10^{-11} cm^{-1} and 1.50×10^{-13} eV. Seventeen software tests covered the archive binding and derivative reconstruction, including rejection of changed bytes, incorrect displacements, nonfinite force/spin values and contradictory backbone labels. These checks establish reproduction of the saved physical evidence; they do not repeat SCF, optimization or Hessian execution, qualify the current GPU runtime, or supply missing conformers.

8.2 Fresh quantum-environment reproduction

A separate CUDA 12.8 compatibility profile retains the recorded molecular-method packages and xTB build while changing fourteen CUDA runtime distributions to satisfy the exact requirements of Torch 2.8.0+cu128. The historical runtime inventory remains separate. All 235 external distributions were installed into a fresh prefix, the project was installed from an archived source checkout, and the dependency check passed. Retrieved package artifacts carry SHA-256 identities; the conda artifact lock and all certificate backing files are retained. The fresh xTB library is byte-identical to the qualified thermodynamic build.

Formaldehyde (singlet) and methyl radical (doublet) were independently optimized and evaluated with the active electronic-energy, native Hessian and projected-partition procedures in the fresh environment. Both passed the normal physical checks and the predeclared original-versus-

fresh comparison bounds: 10^{-5} eV for electronic energy, 10^{-4} eV for free energy, 10^{-3} Å for aligned RMSD and 0.5 cm^{-1} for internal frequencies. A separate replay reconstructed both saved Hessian mode sets and all three rotor-cutoff partitions. Package imports, shared-library locations and the thermodynamic helper interpreter were checked against the fresh prefix; the original environment was not modified. These controls establish reproduction on the tested Linux/CUDA host. They do not establish unrestricted device coverage, eliminate external cluster-reservation configuration or extend polymer conformer validation. The deployment evidence is separate from the earlier cached reviewer-package epoch.

8.3 Relocation of saved native-Hessian partitions

Saved formaldehyde and methyl-radical records were copied with their native Hessians and certificate backing files into a separate checkout. An explicit original-prefix declaration maps the stored certificate paths onto local files only when their hashes and all other certificate metadata agree. The original physical checks remain in force, and the native records are unchanged. With the separately staged, hash-qualified xTB prefix, both projected partitions were recomputed through the CPU entry point. Internal frequencies, symmetry numbers, spin terms and nuclear free energies at all three rotor cutoffs matched their original values exactly. Python reads of the original checkout were blocked; a file-system trace across the process tree recorded no access to files inside it. Ten altered-input controls were rejected, including changed stationarity, gradient, spin, derivative and certificate evidence. These checks establish the tested relocation of saved-Hessian postprocessing; they do not repeat electronic-structure calculations or requalify native GPU Hessian execution.

8.4 Molecular warm-start execution profiles

The software configuration prefers UMA Nitro for gas-phase molecular warm starts on supported hardware and selects the retained standard engine explicitly on V100. The redox-specific current-model/Nitro comparison remains pending; calculations in this study retain their recorded engine identities.[19] Composition, charge and spin define each calculator state; a state change creates a new fixed-state predictor and invalidates cached ASE results. Checkpoint weights, atomic references, loader identity and actual execution mode identify the model profile. Current-model warm starts occupy a separate cache, while archived calculations retain their original model and input identities. Solution energies and nuclear partitions follow the DFT/SMD procedures described above. Numerical comparisons of starting-geometry engines do not establish solution thermochemistry or full polymer conformer convergence.

9 Evidence terminology

Table S12: Terms used in evidence files and their interpretation.

Term	Meaning
Exact design	A specified precursor/product graph and mapped attachment bond; generating recipes can deduplicate to the same design.
Checked or qualified	The checks required for the named field passed; this is not experimental validation.
Native free energy	Active electronic energy plus the projected geometry-level DFT nuclear partition and declared spin partition.
Frozen snapshot	An immutable, hash-verified copy of specified evidence; newer live results require a new snapshot.
Proceed / hold / exclude	Route decisions within the stated evidence scope; hold means required exact-route evidence is incomplete.
Endpoint	One sourced redox-potential comparison, with exact molecular identity, state pair and reference convention.
Pose	One conformational starting geometry; the final optimized contact class is checked separately.

References

[1] N. Harsha Attanayake, Jeffrey A. Kowalski, Katharine V. Greco, Matthew D. Casselman, Jarrod D. Milshtein, Steven J. Chapman, Sean R. Parkin, Fikile R. Brushett, and Susan A. Odom. Tailoring Two-Electron-Donating Phenothiazines To Enable High-Concentration Redox Electrolytes for Use in Nonaqueous Redox Flow Batteries. *Chemistry of Materials*, 31:4353–4363, 2019. doi: 10.1021/acs.chemmater.8b04770.

[2] Rafaela G. Cabral, Gonçalo Viegas, Rita Pacheco, Ana Catarina Sousa, and Maria Paula Robalo. Sustainable Synthesis, Antiproliferative and Acetylcholinesterase Inhibition of 1,4- and 1,2-Naphthoquinone Derivatives. *Molecules*, 28:1232, 2023. doi: 10.3390/molecules28031232.

[3] Mioy T. Huynh, Colin W. Anson, Andrew C. Cavell, Shannon S. Stahl, and Sharon Hammes-Schiffer. Quinone 1 e^- and 2 e^-/2 H^+ Reduction Potentials: Identification and Analysis of Deviations from Systematic Scaling Relationships. *Journal of the American Chemical Society*, 138:15903–15910, 2016. doi: 10.1021/jacs.6b05797.

[4] Huan Wang, Rikard Emanuelsson, Amitava Banerjee, Rajeev Ahuja, Maria Strømme, and Martin Sjödin. Effect of Cycling Ion and Solvent on the Redox Chemistry of Substituted Quinones and Solvent-Induced Breakdown of the Correlation between Redox Potential and Electron-Withdrawing Power of Substituents. *The Journal of Physical Chemistry C*, 124:13609–13617, 2020. doi: 10.1021/acs.jpcc.0c03632.

[5] James B. Gerken and Shannon S. Stahl. High-Potential Electrocatalytic O_2 Reduction with Nitroxyl/NO_x Mediators: Implications for Fuel Cells and Aerobic Oxidation Catalysis. *ACS Central Science*, 1:234–243, 2015. doi: 10.1021/acscentsci.5b00163.

[6] Nicolas Daub, René A. J. Janssen, and Koen H. Hendriks. Imide-Based Multielectron Anolytes as High-Performance Materials in Nonaqueous Redox Flow Batteries. *ACS Applied Energy Materials*, 4:9248–9257, 2021. doi: 10.1021/acsaem.1c01490.

[7] Dukhan Kim, Melanie S. Sanford, Thomas P. Vaid, and Anne J. McNeil. A nonaqueous redox-matched flow battery with charge storage in insoluble polymer beads. *Chemistry—A European Journal*, 28:e202200149, 2022. doi: 10.1002/chem.202200149.

[8] Jan Fulara, Karol Filipkowski, and John P. Maier. Electronic Transition of Ferrocenium: Neon Matrix and CASPT2 Studies. *The Journal of Physical Chemistry C*, 121:10694–10697, 2017. doi: 10.1021/acs.jpcc.6b10391.

[9] Richard J. Bochis, Michael H. Fisher, and Bruce O. Linn. Pyromellitic diimides, processes for their preparation, these compounds in orally administrable form for use in increasing the feed efficiency of a ruminant animal and compositions for this purpose containing them. European patent EP0010440B1, published January 26, 1983, 1983. URL https://patents.google.com/patent/EP0010440B1/en. Example 19, printed p. 10; precursor identity only.

[10] Darren G. Hamilton, Daniel E. Lynch, Karl A. Byriel, Colin H. L. Kennard, and Jeremy K. M. Sanders. N-alkylation of pyromellitic diimide: Solid-state structure of the 1 : 1 N,N′-diethylpyromellitic diimide/2,6-dimethoxynaphthalene cocrystal. *Australian Journal of Chemistry*, 51(5):441–444, 1998. doi: 10.1071/C97214.

[11] Reinhard Sommerlade, Hansjörg Grützmacher, and Souâd Boulmaâz. Process for the selective oxidation of alcohols using readily removable nitroxyl radicals. U.S. Patent 6,441,243 B1, Example A4, 2002. URL https://patents.google.com/patent/US6441243B1/en.

[12] Sofja Tshepelevitsh, Agnes Kütt, Märt Lõkov, Ivari Kaljurand, Jaan Saame, Agnes Heering, Paul G. Plieger, Robert Vianello, and Ivo Leito. On the basicity of organic bases in different media. *European Journal of Organic Chemistry*, 2019:6735–6748, 2019. doi: 10.1002/ejoc.201900956.

[13] University of Tartu, Chair of Analytical Chemistry. Author-maintained basicity collection and recommended acid table. https://analytical.chem.ut.ee/databases/pka-values/. Accessed 22 October 2026. Acid table: https://analytical.chem.ut.ee/HA_UT/pKa_Table_Water_Acetonitrile_DCE_Gas_Phase.pdf.

[14] Josefredo R. Pliego, Jr. Hybrid cluster-continuum method for single-ion solvation free energy in acetonitrile solvent. *The Journal of Physical Chemistry A*, 128:6440–6449, 2024. doi: 10.1021/acs.jpca.4c03593.

[15] D. Barišić, N. Cindro, N. Vidović, N. Bregović, and V. Tomišić. Protonation and anion-binding properties of aromatic sulfonylurea derivatives. *RSC Advances*, 11(39):23992–24000, 2021. doi: 10.1039/D1RA04738H.

[16] Emma N. Cook, Ian M. Courter, Diane A. Dickie, and Charles W. Machan. Controlling product selectivity during dioxygen reduction with Mn complexes using pendent proton donor relays and added base. *Chemical Science*, 15(12):4478–4488, 2024. doi: 10.1039/D3SC02611F.

[17] David J. Heldebrant, Phillip K. Koech, James E. Rainbolt, Feng Zheng, Tricia Smurthwaite, Charles J. Freeman, Merit Oss, and Ivo Leito. Performance of single-component CO2-binding organic liquids (CO2BOLs) for post combustion CO2 capture. *Chemical Engineering Journal*, 171(3):794–800, 2011. doi: 10.1016/j.cej.2011.02.012.

[18] Adrian W. Lange and John M. Herbert. A smooth, nonsingular, and faithful discretization scheme for polarizable continuum models: The switching/gaussian approach. *The Journal of Chemical Physics*, 133(24):244111, 2010. doi: 10.1063/1.3511297.

[19] FAIR Chemistry and Collaborators. Inference using ASE and predictor interface. FAIR Chemistry documentation, 2026. URL https://facebookresearch.github.io/fairchem/ase-calculator/. Accessed October 24, 2026; current model source revision 645af7d.